

Jacob Burkhart



@beanstalksurf

How to Fail at Background Jobs

jacobo.github.com/background_jobs



Delayed::Job



try this instead: (requires postgresql)

[github.com/ryandotsmith/queue classic](https://github.com/ryandotsmith/queue_classic)

The End



DISTILL

distill.engineyard.com

CFP closes Tuesday, April 9

Seeking Better Abstractions for Background Jobs

A photograph of a surfer riding a massive, curling green wave. The surfer is positioned near the base of the wave, facing towards the left. The water is turbulent and white at the bottom of the wave, transitioning to a deep green as it rises.

jacobo.github.com/background_jobs

Rails 4 Queuing



github.com/rails/rails/commit/f9da785d0b1b22317cfca25c15fb555e9016accb

The API

```
class MyJob

  def initialize(account)
    @account = account
  end

  def run
    puts "working on #{@account.name} ..."
  end

end

Rails.queue[:jobs].push(MyJob.new(account))
```

Fail at Serialization

Fail at Serialization

Solve the wrong problem

```
class BodyProxy
  def each
    @body.each{|x| yield x}
    while(email = CleverMailer.emails_to_send.pop)
      email.deliver
    end
  end
end

def call(env)
  status, headers, body = @app.call(env)
  headers["Content-Length"] = body.map(&:bytesize).sum
  [status, headers, BodyProxy.new(body)]
end
```

Will be ~~revisited~~
re-written in 4.1

github.com/rails/rails/pull/9910

github.com/rails/rails/pull/9924

Moving on....



2009

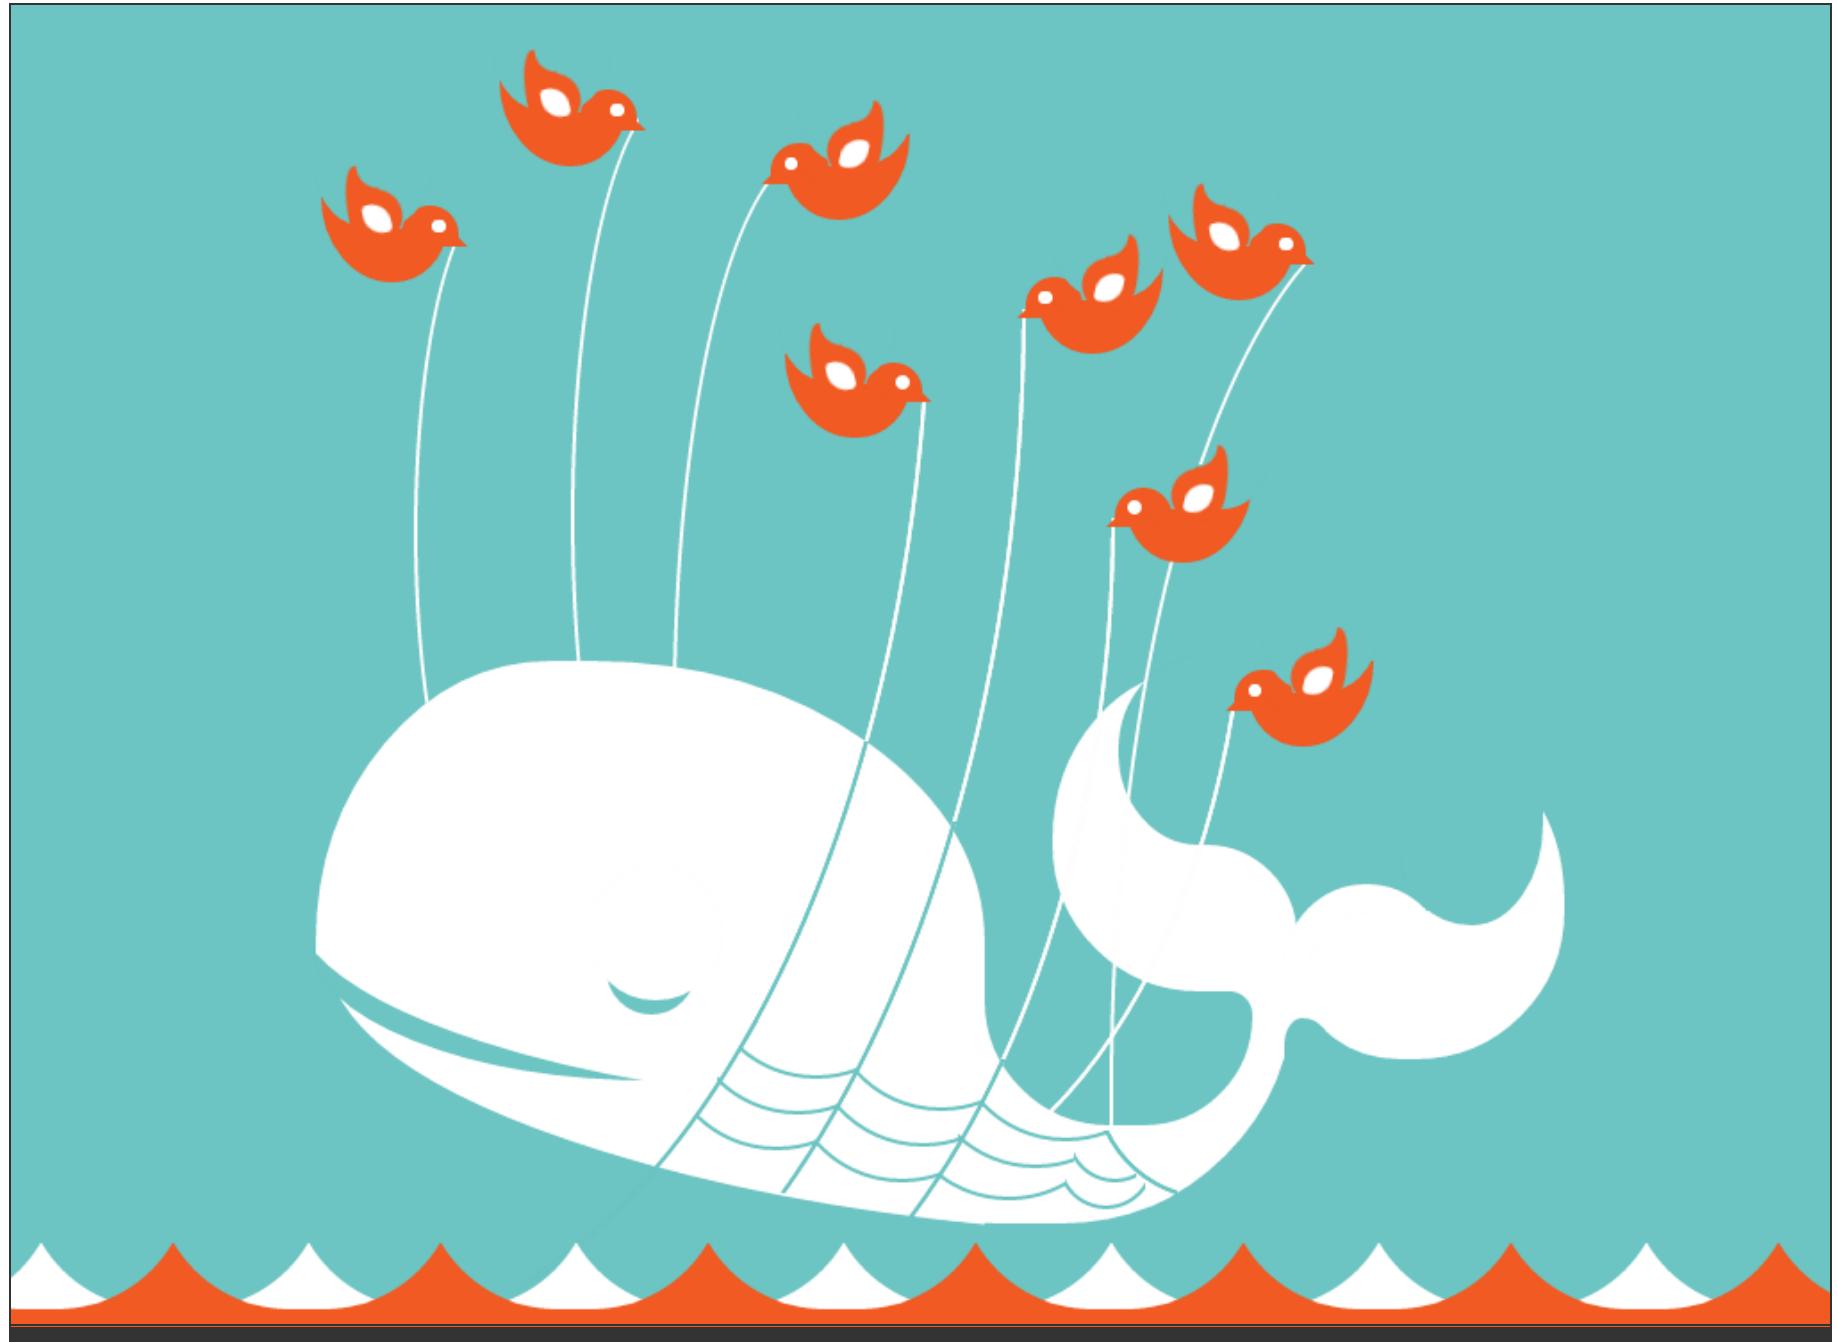


Starling, Workling

```
class FileCopier < Workling::Base
  def copy_case_files(options)
    ...
  end
end
FileCopier.async_copy_case_files(...)

Workling::Remote.dispatcher =
  Workling::Remote::Runners::StarlingRunner.new
```

```
starling -f config/starling.yml
script/workling_client run
```





Installation Documentation Tutorials Services Community Blog

What is RabbitMQ?

- › Robust **messaging** for applications
- › **Easy to use**
- › Runs on all major **operating systems**
- › Supports a huge number of **developer platforms**
- › **Open source and commercially supported**

Get Started

Download

Server and clients for
various operating systems
and languages

Documentation

How to do almost anything
with RabbitMQ

Tutorials

Teach yourself RabbitMQ in
six easy lessons

AMQP

```
connection = AMQP.connect(:host => '127.0.0.1')

channel   = AMQP::Channel.new(connection)
queue     = channel.queue("some.q")
exchange  = channel.default_exchange

queue.subscribe do |payload|
  puts "Received a message: #{payload}"
end

exchange.publish "Hello, world!",
                  :routing_key => queue.name
```

github.com/ruby-amqp/amqp

github.com/ruby-amqp/bunny

 [brontes3d / working](#) watch

forked from [elecnix/working](#)

[Code](#) [Network](#) [Pull Requests 0](#) [Wiki](#)

 tree: [0edc63834a](#) [▼](#) [Files](#) [Commits](#) [Branches 1](#) [+](#)

[working](#) / [lib](#) / [working](#) / [remote](#) / [runners](#) / [+](#)

create initial binding between out queue and exchange

 Simon Horne authored 4 years ago

..

amqp_exchange_runner.rb	4 years ago	create initial binding between out queue and exchange
backgroundjob_runner.rb	4 years ago	refactored pollers, initial work [purzelrakete]
base.rb	4 years ago	more refactoring to accomodate rudeq [purzelrakete]
client_runner.rb	4 years ago	0.4.2, see CHANGES.markdown [purzelrakete]
not_remote_runner.rb	4 years ago	added loads of missing method and class documentation
spawn_runner.rb	4 years ago	fix another minor typo in the spawn runner. [Brian Coo
starling_runner.rb	4 years ago	0.4.1 see changelog for details [purzelrakete]

Little Bug



ActiveRecord:: RecordNotFound

```
class Device < ActiveRecord::Base
  after_create do |d|
    BackgroundJob.async_run(d.id)
  end
end

class BackgroundJob
  def run(device_id)
    #sleep 1 ?
    Device.find(device_id)
  end
end
```

Hack ActiveRecord

```
class Device < ActiveRecord::Base
  after_create do |d|
    d.commit_callback do
      BackgroundJob.async_run(d.id)
    end
  end
end

...
def commit_callback
  self.connection.instance_eval do
    class << self
      alias commit_db_transaction_original commit_db_transac
      ...
    end
  end
end
```

github.com/brontes3d/commit_callback

More Bugs



Fundamental Flaw



Do it Yourself

```
class FileCopier < AmqpListener::Listener
  subscribes_to :case_file_copy_requests

  def handle(options)
    ...
  end
end

FileCopier.notify(...)
```

```
script/amqp_listener run
```

github.com/brontes3d/amqp_listener

Moment of Reflection

A photograph of a sunset over a calm ocean. The sky is a gradient from deep blue at the top to a bright orange and yellow near the horizon. Silhouettes of tropical trees and leaves frame the scene. The ocean reflects the warm colors of the sunset.

Trains and Resque



Boot an EC2 Server

```
class InstanceProvision

  def self.perform(instance_id)
    instance = Instance.find(instance_id)

    fog = Fog::Compute.new(...)
    server = fog.servers.create(...)
    instance.amazon_id = server.id

    while(!server.ready?)
      sleep 1
      server.reload
    end

    instance.attach_ip!
    ...
  end
end
```

Extractable?

```
class InstanceProvision

  def self.perform(aws_creds, create_params, callback_url)
    fog = Fog::Compute.new(aws_creds)
    server = fog.servers.create(create_params)
    API.post(callback_url, :instance_amazon_id => server.id)

    while(!server.ready?)
      sleep 1
      server.reload
    end

    ip = fog.ips.create!
    ip.server = server
    API.post(callback_url, :attached_ip_amazon_id => ip.id)
    ...
  end
end
```

Generalizable?

```
class MethodCalling
  def self.perform(class_str, method, id, *args)
    model_class = Object.const_get(class_str)
    model = model_class.find(id)
    model.send(method, *args)
  end
end

class Instance
  def provision
    Resque.enqueue(MethodCalling, Instance, :provision!, id)
  end

  def provision!
    #actually do it
  end
end
```

Async

```
require 'async'
require 'async/resque'
Async.backend = Async::ResqueBackend

class Instance < ActiveRecord::Base
  def provision(*args)
    Async.run{ provision_now(*args) }
  end
  def provision_now(*args)
    #actually do it
  end
end
```

github.com/engineyard/async

Application Instances

 Application High CPU Medium (32 bit)	i-f3b19dc1 SSH / HTTP	Logs: Base / Custom Performance graphs	 Terminate
Waiting - about 6 hours ago			
Waiting - about 6 hours ago			
Waiting - about 6 hours ago			
 Application Master High CPU Medium (32 bit)	i-8fb19dbd SSH / HTTP	Logs: Base / Custom Performance graphs	
Waiting - about 6 hours ago			
Waiting - about 6 hours ago			
Waiting - about 6 hours ago			

Database Instances

 Master Database High CPU Medium (32 bit)	i-85b19db7 SSH	Logs: Base / Custom Performance graphs
Waiting - about 6 hours ago		
Waiting - about 6 hours ago		
Waiting - about 6 hours ago		

More Options

[Edit Environment](#)
Make Stack Changes

[Crontabs](#)
Add or Edit a Cron Job

46 of 157 Workers Working

The list below contains all workers which are currently running a job.

	Where	Queue	Processing
1	ip-25-100-68-98:11991	INSTANCE_PROVISIONING	InstanceProvisioning about 14 hours ago
2	ip-25-5-66-63:18066	INSTANCE_PROVISIONING	InstanceProvisioning about 12 hours ago
3	domU-22-31-38-5F-01-E1:6053	INSTANCE_PROVISIONING	InstanceProvisioning about 10 hours ago
4	domU-22-31-38-5E-51-31:6379	INSTANCE_PROVISIONING	InstanceProvisioning about 10 hours ago
5	ip-25-89-63-160:31812	INSTANCE_PROVISIONING	InstanceProvisioning about 8 hours ago
6	ip-25-126-111-88:19063	INSTANCE_PROVISIONING	InstanceProvisioning about 7 hours ago
7	domU-22-31-38-5E-51-31:6449	INSTANCE_PROVISIONING	InstanceProvisioning about 7 hours ago
8	ip-25-80-39-43:24306	INSTANCE_PROVISIONING	InstanceProvisioning about 7 hours ago
9	domU-22-31-38-5E-51-31:6760	INSTANCE_PROVISIONING	InstanceProvisioning about 4 hours ago
10	domU-22-31-38-5E-51-31:6670	INSTANCE_PROVISIONING	InstanceProvisioning about 2 hours ago
11	domU-22-31-38-5E-51-31:6370	INSTANCE_PROVISIONING	InstanceProvisioning about 2 hours ago
12	ip-25-80-39-43:24268	INSTANCE_PROVISIONING	InstanceProvisioning about 2 hours ago
13	ip-25-126-111-88:19161	INSTANCE_PROVISIONING	InstanceProvisioning about 2 hours ago
14	ip-25-89-63-160:31872	INSTANCE_PROVISIONING	InstanceProvisioning about 2 hours ago

Reliability

Reliable Queue vs Reliable Job

Retry (raise, crash, hang). Time-out

Monitor / Maintain N workers

Graceful Restart

Did it run? Did it fail? How? Where?

Make a Resque Plugin

```
class InstanceProvision
  extend Resque::Plugins::JobTracking

  def self.track(instance_id, opts)
    i = Instance.find(instance_id)
    ["Account:#{i.account_id}",
     "Instance:#{instance_id}"]
  end

  def self.perform(instance_id, opts)
    #do stuff
  end
end
```

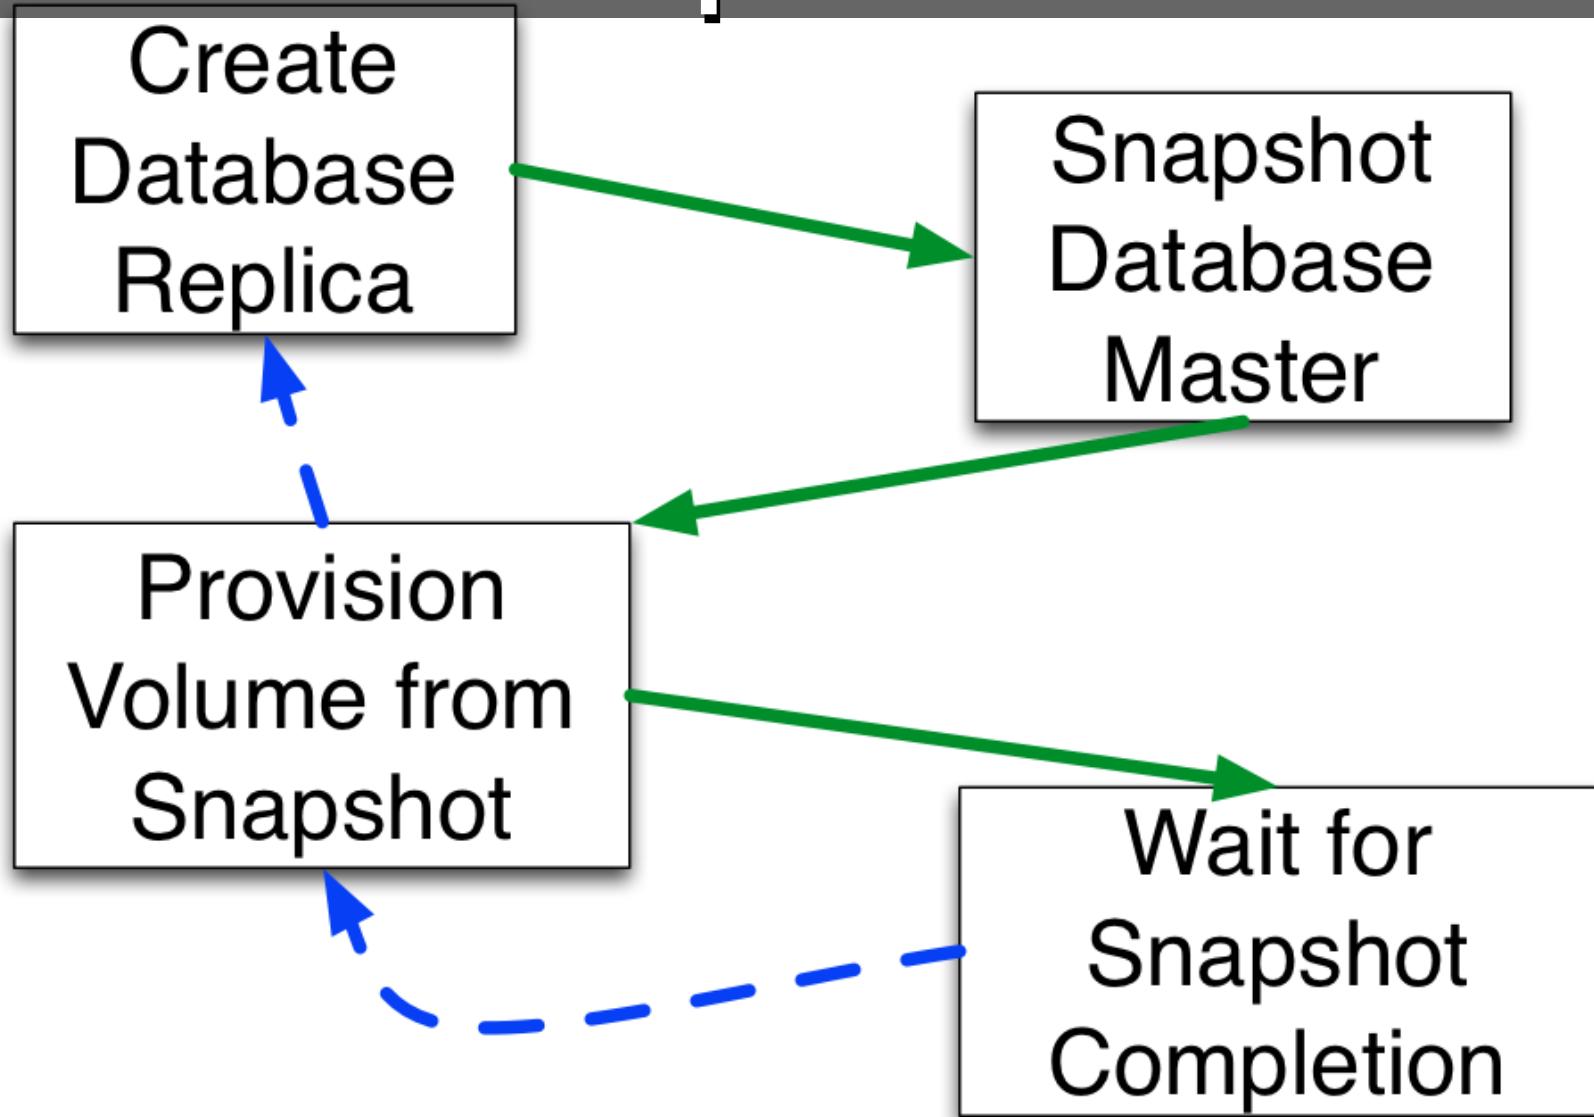
github.com/engineyard/resque-job-tracking

Helps a little?

```
InstanceProvision.pending_jobs("Instance:532")
```

```
InstanceProvision.failed_jobs("Account:121")
```

Jobs Dependencies



Make another

```
class Sandwich
  extend Resque::Plugins::Delegation

  def self.steps(tomato_color, cheese_please, cheesemaker)
    step "fetch a", :tomato do
      depend_on(Tomato, tomato_color)
    end
    step "slice the ", :tomato, " and make", :tomato_slices
      tomato.split(",")
    end
    step "fetch the", :cheese_slices do
      if cheese_please
        depend_on(Cheese, cheesemaker)
      ...
    end
  end
end
```

github.com/engineyard/resque-delegation

Desperation?

```
class ResqueJob < ActiveRecord::Base  
end
```

```
class AbstractJob  
  
  def self.store_meta(meta)  
    meta_id = meta["meta_id"]  
    ResqueJob.create! (meta.slice(  
      "meta_id", "job_class",  
      "enqueued_at", "started_at", "finished_at"))  
    super(meta)  
  end  
  
  ...
```

Maybe we just need better logging?

<http://youtu.be/NpTT30wLL-w>

Model Intent



Data belongs in a database (not redis)

```
class InstanceProvision < ActiveRecord::Base
  belongs_to :instance

  def run(instance_id)
    fog = Fog::Compute.new(...)
    server = fog.servers.create(...)
    instance.amazon_id = server.id

    while(!server.ready?)
      sleep 1
      server.reload
    end

    instance.attach_ip!
```

requested_at

started_at

finished_at

state

Idempotent

```
class InstanceProvision < ActiveRecord::Base
  belongs_to :instance

  def run(instance_id)
    with_lock do
      if self.state == "running"
        raise
      else
        self.state = "running"
        self.started_at = Time.now
        save!
      end
    end
  ...
  ...
```

Non-Resque specific job-middleware

```
procedure = Viaduct::Builder.new do
  use Instrumentation
  use ProvisionNotification
  use ProvisionInstance
end
Viaduct::Runner.new.run(procedure, data)
```

github.com/slack/viaduct

Moment of Reflection





3 Ingredients

Work Loop

Monitor & Restart

Queue

Loop



Resque Event Loop

```
def work(interval = 5, &block)
  loop do
    run_hook :before_fork, job

    if @child = fork
      procline "Forked #{@child} at #{Time.now.to_i}"
      Process.wait
    else
      procline "Processing #{job.queue} since #{Time.now.to_i}"
      perform(job, &block)
      exit! unless @cant_fork
    end
  end
end
```

github.com/defunkt/resque/blob/master/lib/resque/worker.rb

EventMachine

```
void EventMachine_t::Run()
//Epoll and Kqueue stuff...
...
while (true) {
    _UpdateTime();
    _RunTimers();

    _AddNewDescriptors();
    _ModifyDescriptors();

    _RunOnce();
    if (bTerminateSignalReceived)
        break;
}
```

github.com/eventmachine/eventmachine/blob/master/ext/em.cpp

EM.next_tick

```
require 'eventmachine'  
EM.run {  
    EM.start_server(host, port, self)  
}  
  
EM.next_tick{ puts "do something" }
```

Celluloid / Sidekiq

```
class Sidekiq::Manager
  include Celluloid
  ????

class Sidekiq::Fetcher
  include Celluloid
  ???

class Sidekiq::Processor
  include Celluloid
  ????
```

sidekiq.org

github.com/celluloid/celluloid

Monitor & Restart



Resque Signals

QUIT - Wait for child to finish then exit

TERM / INT - Immediately kill child, exit

USR1 - Immediately kill child, don't exit

USR2 - Don't start to process new jobs

CONT - Start to process new jobs again
after a USR2

Monitor with God

```
5.times do |n|
  God.watch do |w|
    w.name      = "resque-#{num}"
    w.group     = 'resque'
    w.interval  = 30.seconds
    w.log       = "#{app_root}/log/worker.#{num}.log"
    w.dir       = app_root
    w.env       = {
      "GOD_WATCH"  => w.name,
      "QUEUE"       => '*'
    }
    w.start     = "bundle exec rake --trace resque:work"
  ...
}
```

godrb.com

Or Daemons

```
require 'daemons'

options = {
  :app_name => "worker",
  :log_output => true,
  :backtrace => true,
  :dir_mode => :normal,
  :dir => File.expand_path('../tmp/pids', __FILE__),
  :log_dir => File.expand_path('../log', __FILE__),
  :multiple => true,
  :monitor => true
}

Daemons.run(File.expand_path('../worker', __FILE__), option
```

daemons.rubyforge.org

Or Don't?



torquebox.org



Queue



Thread-safe Queue

```
module BundlerApi
  class ConsumerPool

    def initialize
      @queue = Queue.new
    end

    def enq(job)
      @queue.enq(job)
    end

    ...
    job = @queue.deq
    job.run
  end
end
```

github.com/rubygems/bundler-api/blob/master/lib/bundler_api/update/consumer_pool.rb

DRB objects

```
class MultiHeadedGreekMonster
  class ServiceManager
    def initialize(progress)
      @things = []
    end
    def give(thing)
      @things.unshift(thing)
    end
    def take
      @things && @things.pop
    end
  ...
  DRb.start_service "druby://localhost:#{@on_port}",
    ServiceManager.new(@progress)
```

https://github.com/engineyard/multi_headed_greek_monster/blob/master/

Qu Push/Pop Redis

```
payload.id = SimpleUUID::UUID.new.to_guid
redis.set("job:#{payload.id}",
  encode('klass' => payload.class.to_s,
         'args' => payload.args))
redis.rpush("queue:#{payload.queue}", payload.id)
redis.sadd('queues', payload.queue)

...
redis.lpop(queue)
```

github.com/bkeepers/qu/blob/master/lib/qu/backend/redis.rb

Qu Push/Pop Mongo

```
payload.id = SimpleUUID::UUID.new.to_guid
payload.id = BSON::ObjectId.new
jobs(payload.queue).insert({
  '_id' => payload.id,
  'klass' => payload(klass.to_s, 'args' => payload.args))
self[:queues].update({'name' => payload.queue},
  {'name' => payload.queue}, 'upsert' => true)

...
if doc = jobs(queue).find_and_modify(:remove => true)
  doc['id'] = doc.delete('_id')
  return Payload.new(doc)
end
```

github.com/bkeepers/qu/blob/master/lib/qu/backend/mongo.rb

All Together



script/invoice_worker

```
#!/usr/bin/env ruby
require File.expand_path('../config/environment',
__FILE__)
TrapLoop.start do
  InvoiceProcessingTask.process_invoices!
end
```

script/invoice_runner start

script/invoice_runner stop

```
require 'daemons'
Daemons.run(File.expand_path('../invoice_worker', __FILE__))
```

```
class TrapLoop
  trap('TERM') { stop! }
  trap('INT') { stop! }
  trap('SIGTERM') { stop! }

  def self.stop!
    @loop = false
  end

  def self.safe_exit_point!
    if @started && !@loop
      raise Interrupt
    end
  end

  def self.start(&block)
    @started = true
    @loop = true
    while(@loop) do
      yield
    end
  end
end
```

```
class InvoiceProcessingTask
  def self.enq_task(task_id, invoice_id)
    REDIS.rpush("tasks:#{invoice_id}", task_id)
    REDIS.rpush("invoices", invoice_id)
  end

  def self.process_invoices!
    while invoice_id = REDIS.lpop("invoices")
      process_tasks!(invoice_id)
      TrapLoop.safe_exit_point!
    end
  end

  def self.process_tasks!(invoice_id)
    while task_id = REDIS.lpop("tasks:#{invoice_id}")
      if task = InvoiceProcessingTask.find_by_id(task_id)
        task.process!
        TrapLoop.safe_exit_point!
      end
    end
  end
end
```

Conclusions

Abstraction Awareness

Model important jobs in
your domain

Contribute to Resque

Questions?

A photograph of a group of people, likely actors or performers, with their faces and hair completely covered in a dark, mottled paint. They are looking directly at the camera with serious expressions. The lighting is dramatic, with strong highlights on their faces against a dark background.

jacobo.github.com/background_jobs

@beanstalksurf

jacobo.github.com/talks